

Operating Systems

Chapter 4 I/O and Driver

Tien Pham Van, Dr. rer. nat.

Compiled with reference to other presentations

- Communication between an information processing system (computer) and peripherals, or another processing system
- It may refer to input/output operations for data exchange

- Two fundamental operations performed by a computing system
 - Processing of data (implement an algorithm)
 - Perform I/O (move data into and out of the system)
- Large diversity in I/O devices, capabilities and performance
 - Common Categories: storage, transmission and user-interface devices.
 - Examples: display, keyboard, network, hard disks, tape drives, CDROM ...
- OS Goal is to
 1. provide simple and consistent interface to user
 2. optimize I/O use for maximum concurrency
- Mechanisms used by OS
 - device drivers provide standard interface to I/O devices

Categories of I/O Devices

- Human readable
 - Used to communicate with the user
 - Printers
 - Video display terminals
 - Display
 - Keyboard
 - Mouse

Categories of I/O Devices

- Machine readable
 - Used to communicate with electronic equipment
 - Disk and tape drives
 - Sensors
 - Controllers
 - Actuators

Categories of I/O Devices

- Communication
 - Used to communicate with remote devices
 - Digital line drivers
 - Modems

Differences in I/O Devices

- Data rate
 - May be differences of several orders of magnitude between the data transfer rates




Figure 11.1 Typical I/O Device Data Rates

Differences in I/O Devices

- Application
 - Disk used to store files requires file management software
 - Disk used to store virtual memory pages needs special hardware and software to support it
 - Terminal used by system administrator may have a higher priority

Differences in I/O Devices

- Complexity of control
- Unit of transfer
 - Data may be transferred as a stream of bytes for a terminal or in larger blocks for a disk
- Data representation
 - Encoding schemes
- Error conditions
 - Devices respond to errors differently

Performing I/O

- Programmed I/O
 - Process is busy-waiting for the operation to complete
- Interrupt-driven I/O
 - I/O command is issued
 - Processor continues executing instructions
 - I/O module sends an interrupt when done

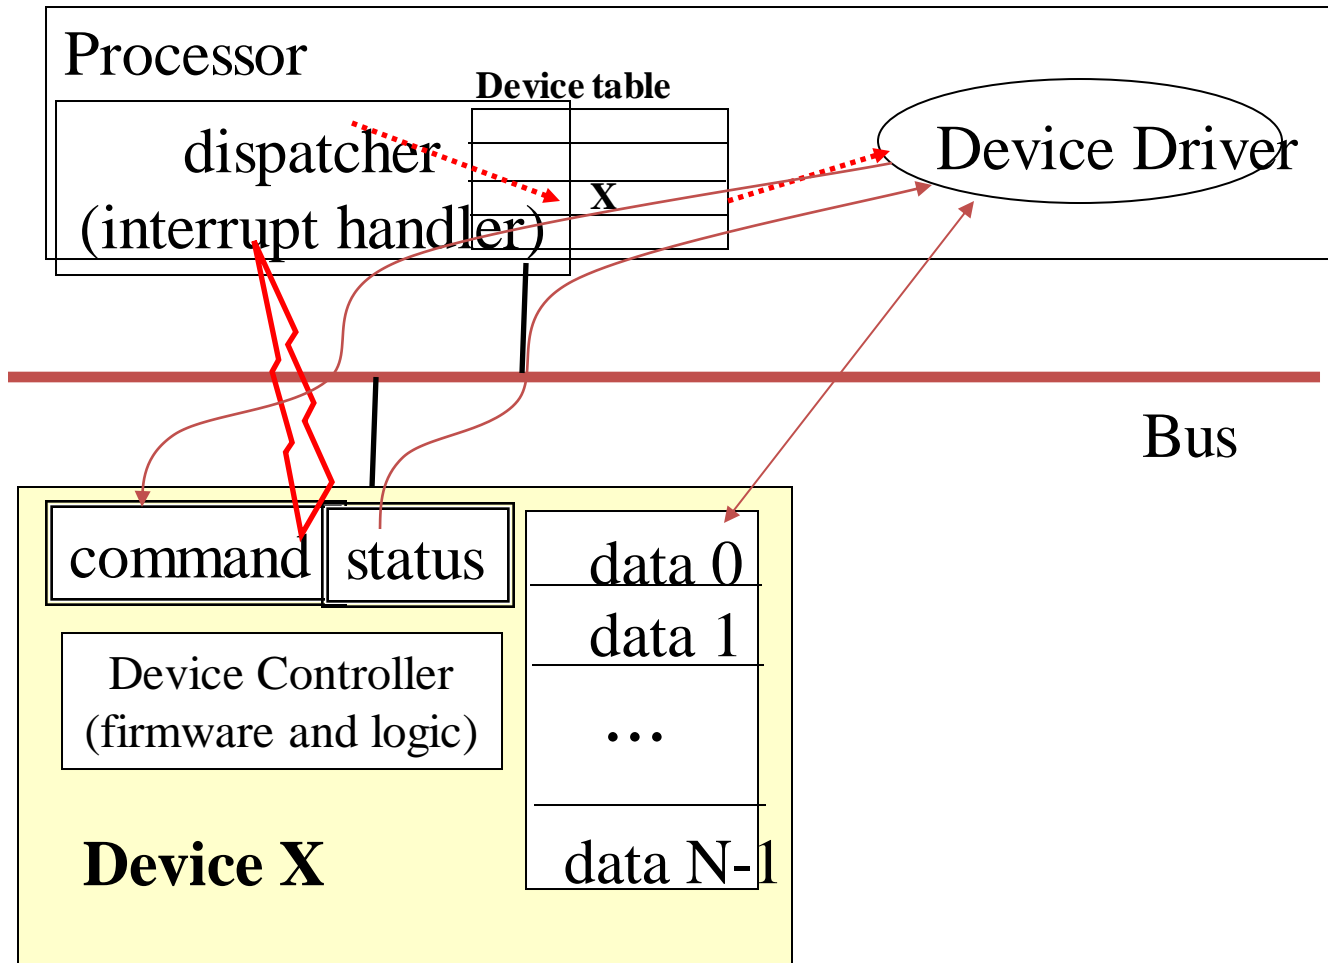
Performing I/O

- Direct Memory Access (DMA)
 - DMA module controls exchange of data between main memory and the I/O device
 - Processor interrupted only after entire block has been transferred

Table 11.1 I/O Techniques

	No Interrupts	Use of Interrupts
I/O-to-memory transfer through processor	Programmed I/O	Interrupt-driven I/O
Direct I/O-to-memory transfer		Direct memory access (DMA)

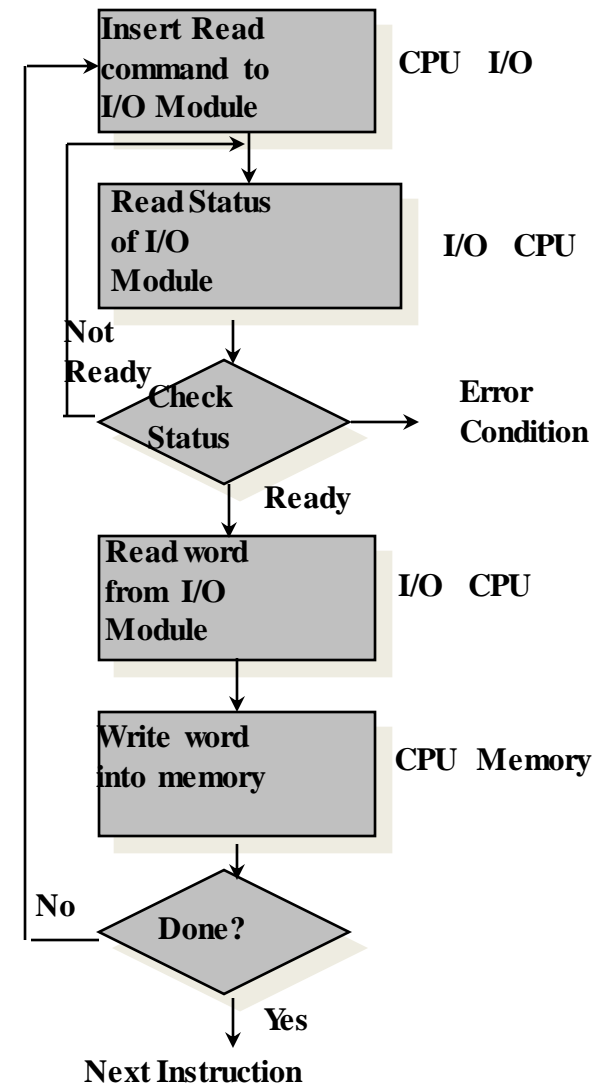
I/O and Devices



- Common concepts
 - *Port* (device-machine communication point)
 - *Bus* (daisy chain or shared communication channel)
 - *Controller* (host adapter)
- Common Techniques
 1. Direct I/O with polling, aka *Programmed I/O*: Processor does all the work. Poll for results.
 2. Interrupt Driven I/O: Device notifies CPU when I/O operation complete
 3. Memory Mapped I/O: rather than reading/writing to controller registers the device is mapped into the OS memory space. increased efficiency
 4. Direct memory access (DMA): DMA controller read and write directly to memory, freeing the CPU to do other things. CPU notified when DMA complete

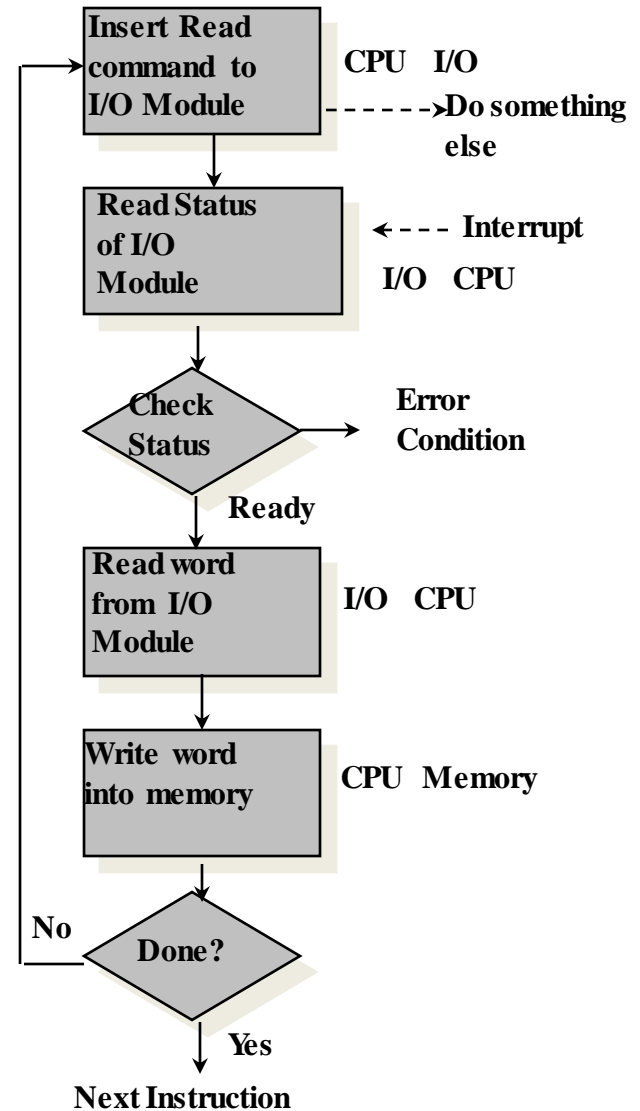
Programmed I/O

- *Busy-wait* cycle to wait for I/O from device
- Poll at select times: periodic, entering/leaving kernel etc.
 - Determines state of device: **command-ready, busy, error**
- Processor transfer data to/from device.
- Read/write directly to status and command registers
- Processor polls device for status
- Consumes a lot of processor time because every word read or written passes through the processor



Interrupt-Driven I/O

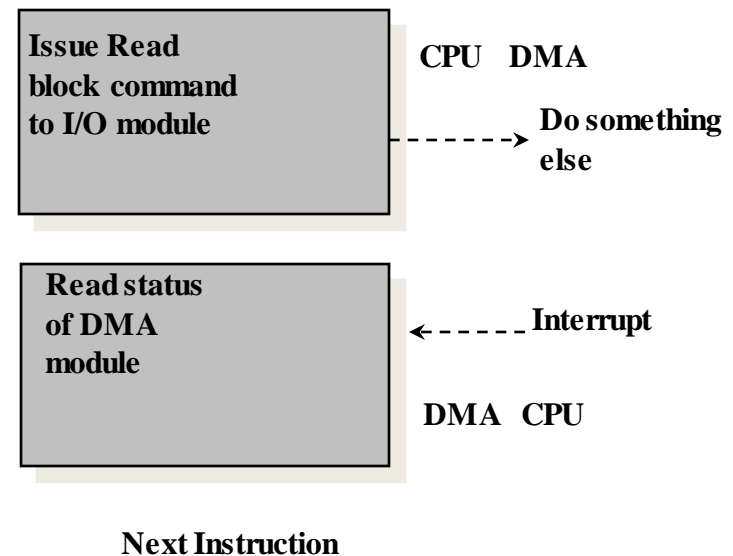
- Similar to direct I/O but processor not required to poll device.
- Interrupt asserted to notify processor of a change in status
- *Interrupt handler* receives interrupts
- *Maskable* to ignore or delay some interrupts
- *Interrupt vector* to dispatch interrupt to correct handler
 - Based on priority
 - Some unmaskable



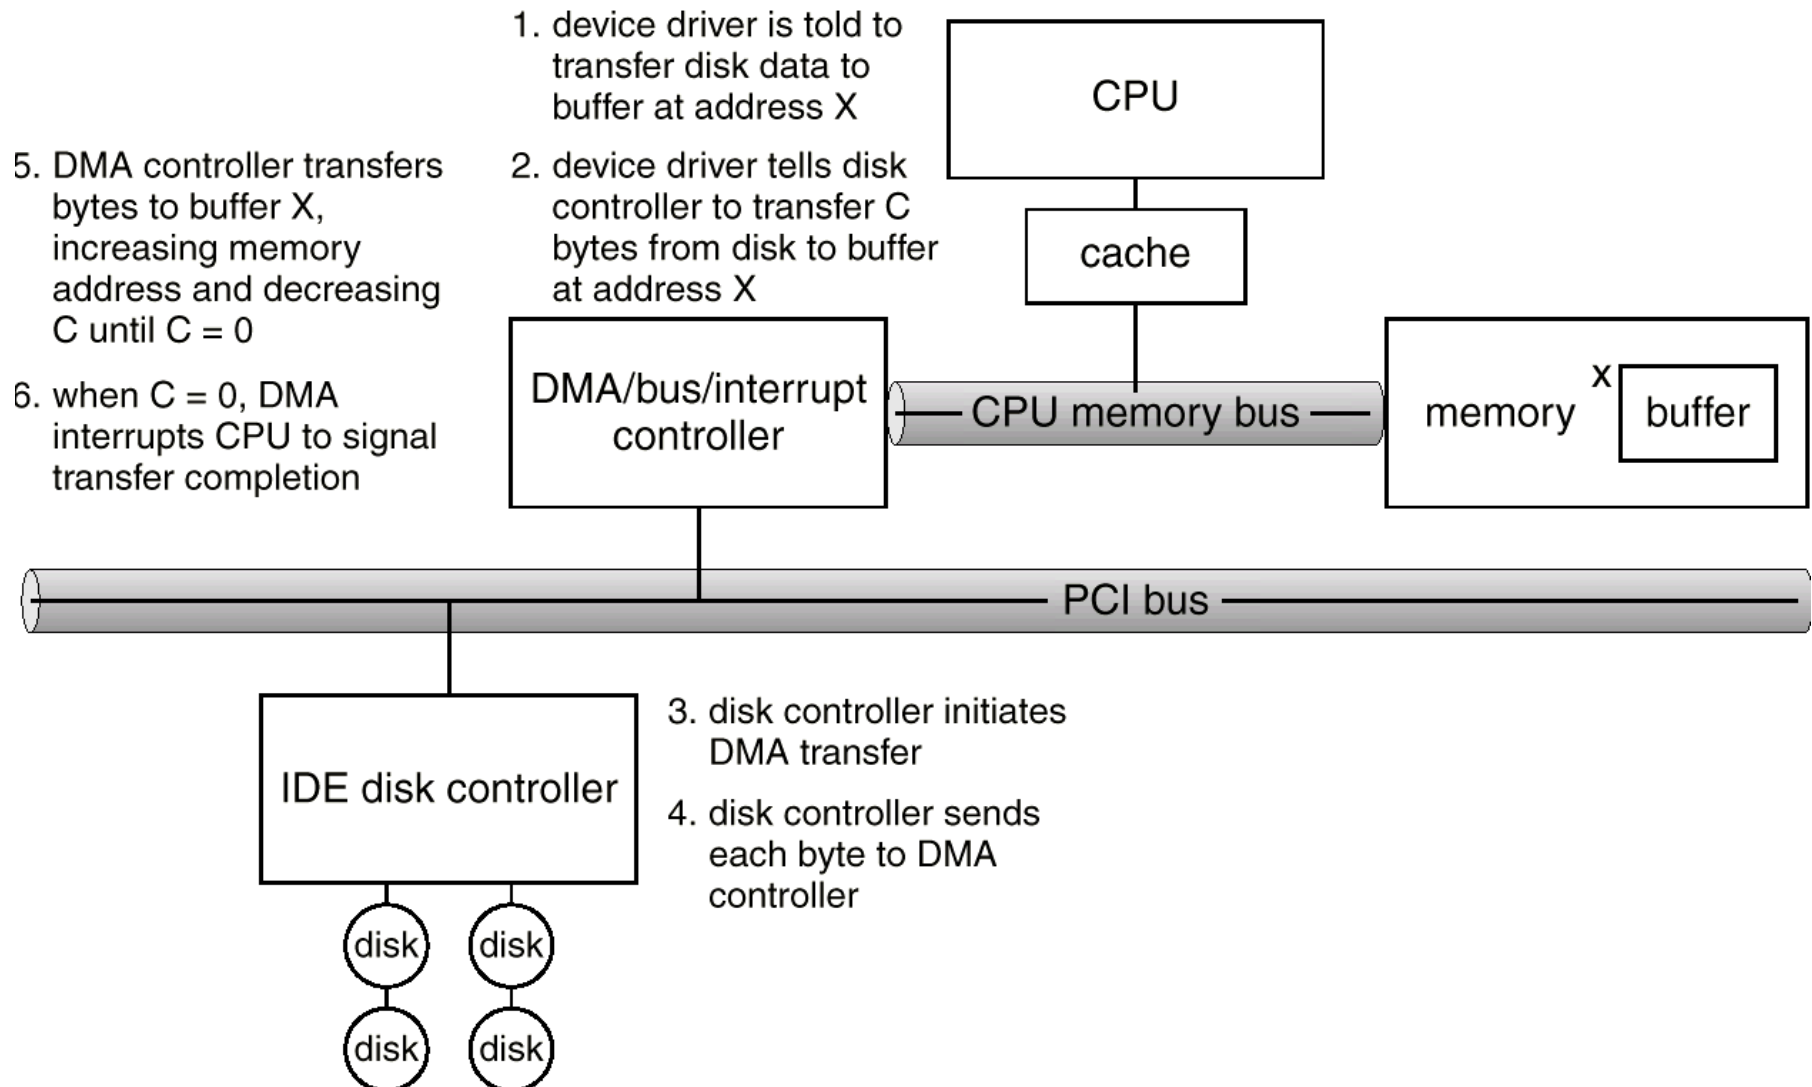
Interrupt mechanism also used for exceptions

Direct Memory Access (DMA)


- I/O exchanges to memory
 - Processor grants I/O module authority to read from or write to memory
 - Relieves the processor from the task
 - Processor is free to do other things
- An interrupt is sent when the task is complete
- The processor is only involved at the beginning and end of the transfer
- Used to avoid programmed I/O for large data movement
- Requires DMA controller



Six steps to perform DMA transfer



Programmed I/O (1)



Steps in printing a string

Programmed I/O (2)

```
copy_from_user(buffer, p, count);  
for (i = 0; i < count; i++) {  
    while (*printer_status_reg != READY) ;  
    *printer_data_register = p[i];  
}  
return_to_user();
```

/* p is the kernel bufer */
/* loop on every character */
/* loop until ready */
/* output one character */

Writing a string to the printer using programmed I/O

Interrupt-Driven I/O

```
copy_from_user(buffer, p, count);  
enable_interrupts();  
while (*printer_status_reg != READY) ;  
*printer_data_register = p[0];  
scheduler();
```

(a)

```
if (count == 0) {  
    unblock_user();  
} else {  
    *printer_data_register = p[i];  
    count = count - 1;  
    i = i + 1;  
}  
acknowledge_interrupt();  
return_from_interrupt();
```

(b)

- Writing a string to the printer using interrupt-driven I/O
 - Code executed when print system call is made
 - Interrupt service procedure

I/O Using DMA

```
copy_from_user(buffer, p, count);  
set_up_DMA_controller( );  
scheduler( );
```

(a)

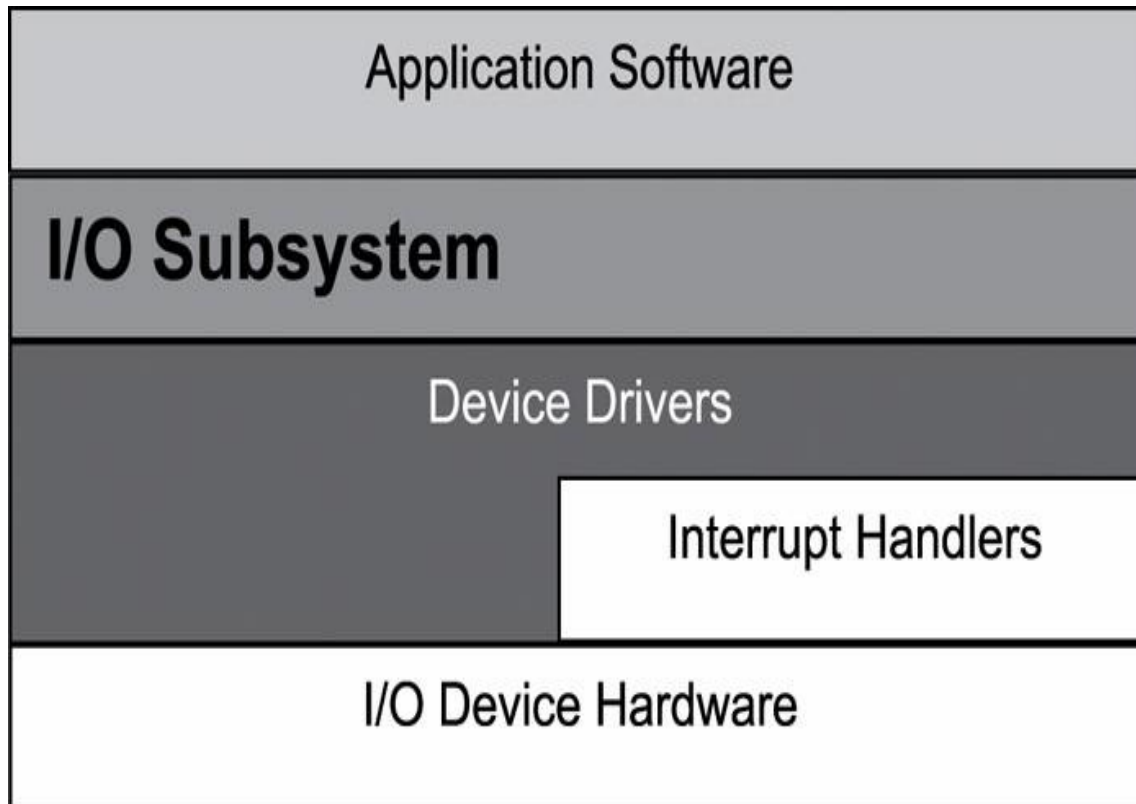
```
acknowledge_interrupt( );  
unblock_user( );  
return_from_interrupt( );
```

(b)

- Printing a string using DMA
 - code executed when the print system call is made
 - interrupt service procedure

- https://www.youtube.com/watch?v=DYGrqNBWymw&ab_channel=GnanaTeja

I/O Subsystem and the Layered Design



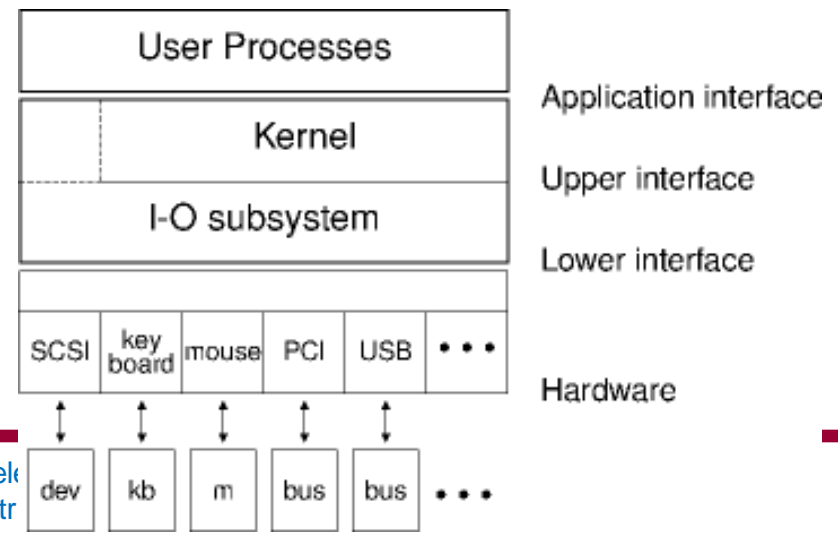
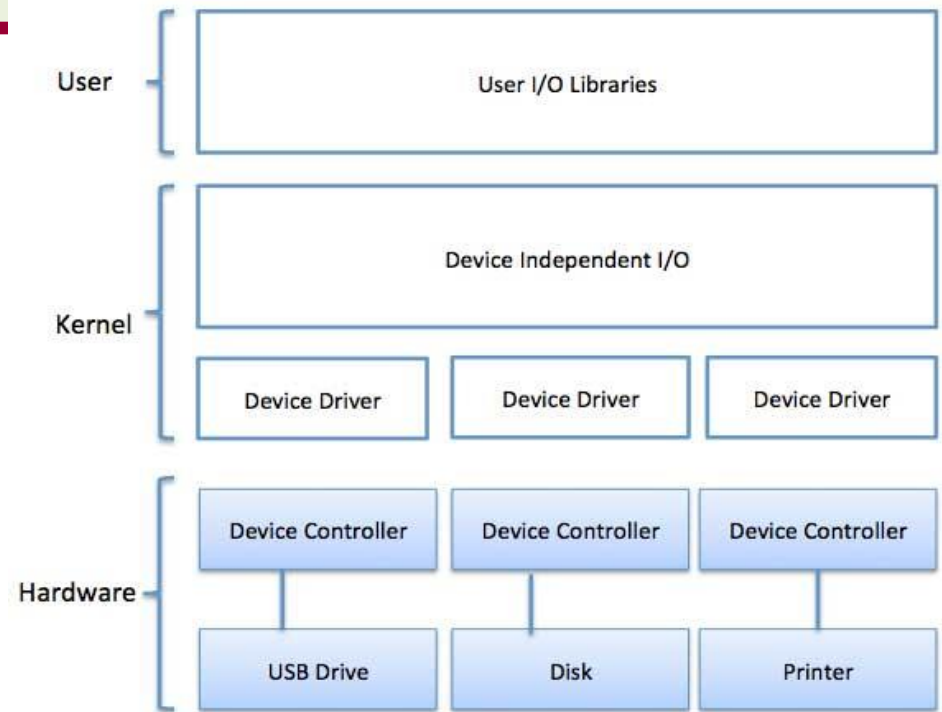
Generic



Specific Details

IO Subsystem

- ❑ Each I/O device driver can provide a driver-specific set of I/O API to applications
- ❑ Purpose of IO subsystem in embedded systems
 - To hide device-specific information from the kernel and from application developer
 - To provide a uniform access method to the peripheral IO devices



IO Subsystem

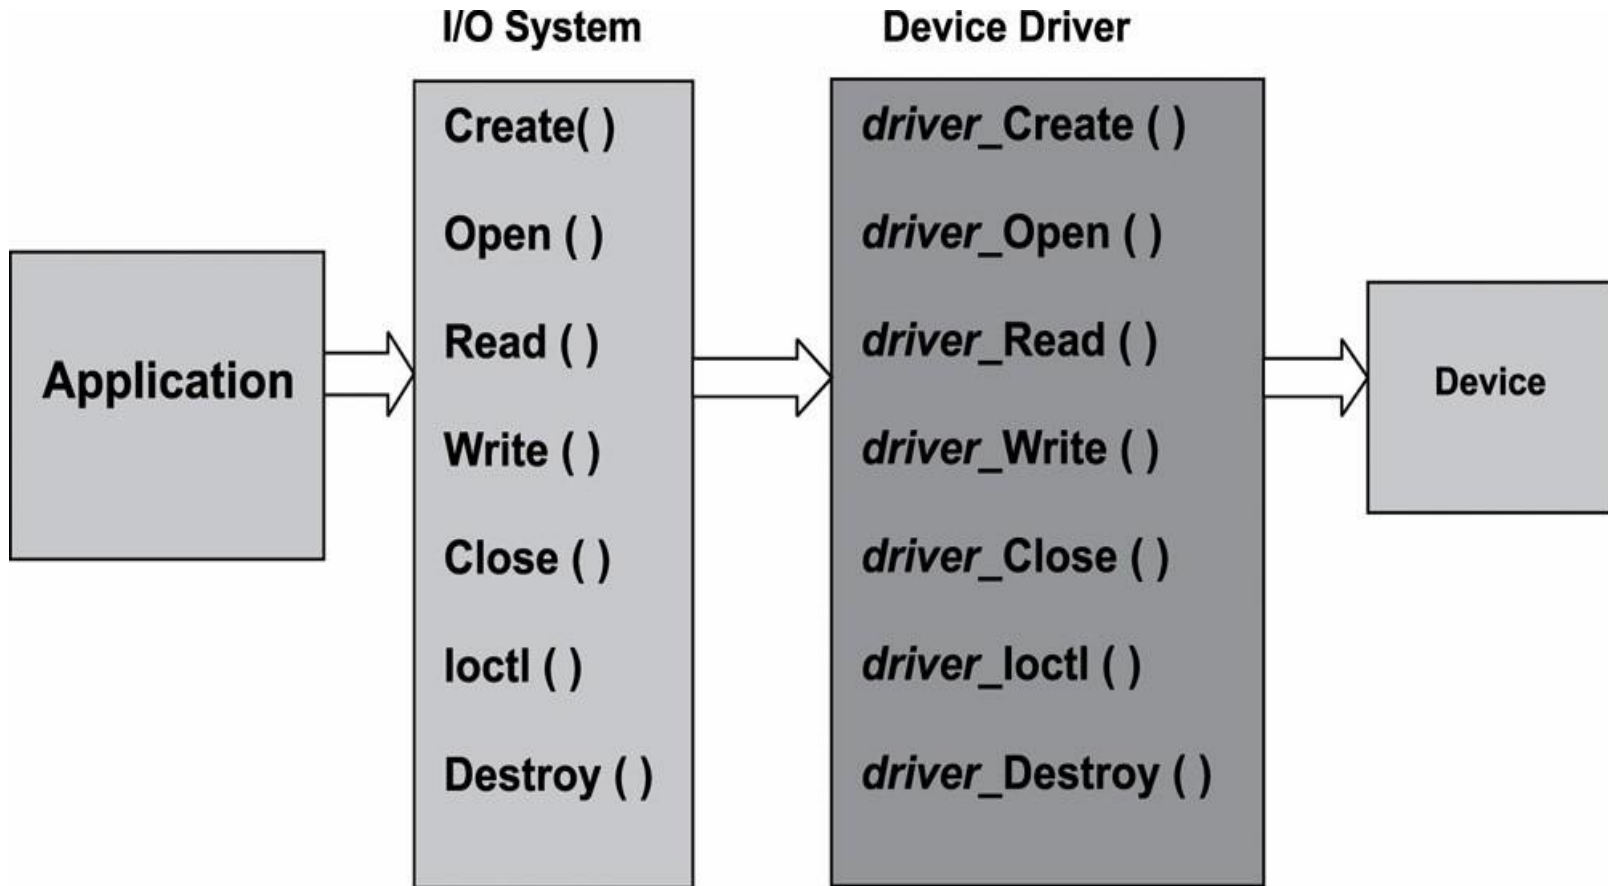
- IO subsystem defines *a standard set of functions* for IO operations
 - To hide device peculiarities from applications

- All IO device drivers conform to and support this function set
 - To provide uniform IO to applications across a wide spectrum of IO devices of varying types

IO Functions

Function	Description
Create	Creates a virtual instance of an I/O device
Destroy	Deletes a virtual instance of an I/O device
Open	Prepares an I/O device for use.
Close	Communicates to the device that its services are no longer required, which typically initiates device-specific cleanup operations.
Read	Reads data from an I/O device
Write	Writes data into an I/O device
ioctl	Issues control commands to the I/O device (I/O control)

IO Function Mapping



Uniform IO Driver Table

Driver Table

	Create	Destroy	Open	Close	Read	Write	ioctl
"fei"			⋮				
"tty"	⋮		⋮				

```
int tty_Create()
{
}

```

```
int fei_Open()
{
}

```

Device tree

- Devicetree is a data structure describing the hardware components of a particular computer so that the kernel can use and manage those components, including the CPUs, memory, buses and peripherals
- It is a tree of named nodes and properties. Nodes contain properties and child nodes, while properties are name–value pairs

Block and Character Devices

- Block devices include disk drives
 - Commands include read, write, seek
 - Raw I/O or file-system access
 - Memory-mapped file access possible
- Character devices include keyboards, mice, serial ports
 - Commands include `get`, `put`
 - Libraries layered on top allow line editing

- *Scheduling*
 - Efficiency, fairness, prioritized access
- *Buffering* - store data in memory while transferring between devices
 - To cope with device speed mismatch, transfer size mismatch and to maintain “copy semantics”
- *Caching* - fast memory holding copy of data
 - Always just a copy, Key to performance
- *Spooling* - hold output for a device
 - Used when device can serve only one request at a time, i.e., Printing
- *Device reservation* - provides exclusive access to a device
 - System calls for allocation and deallocation
 - Watch out for deadlock

Error Handling

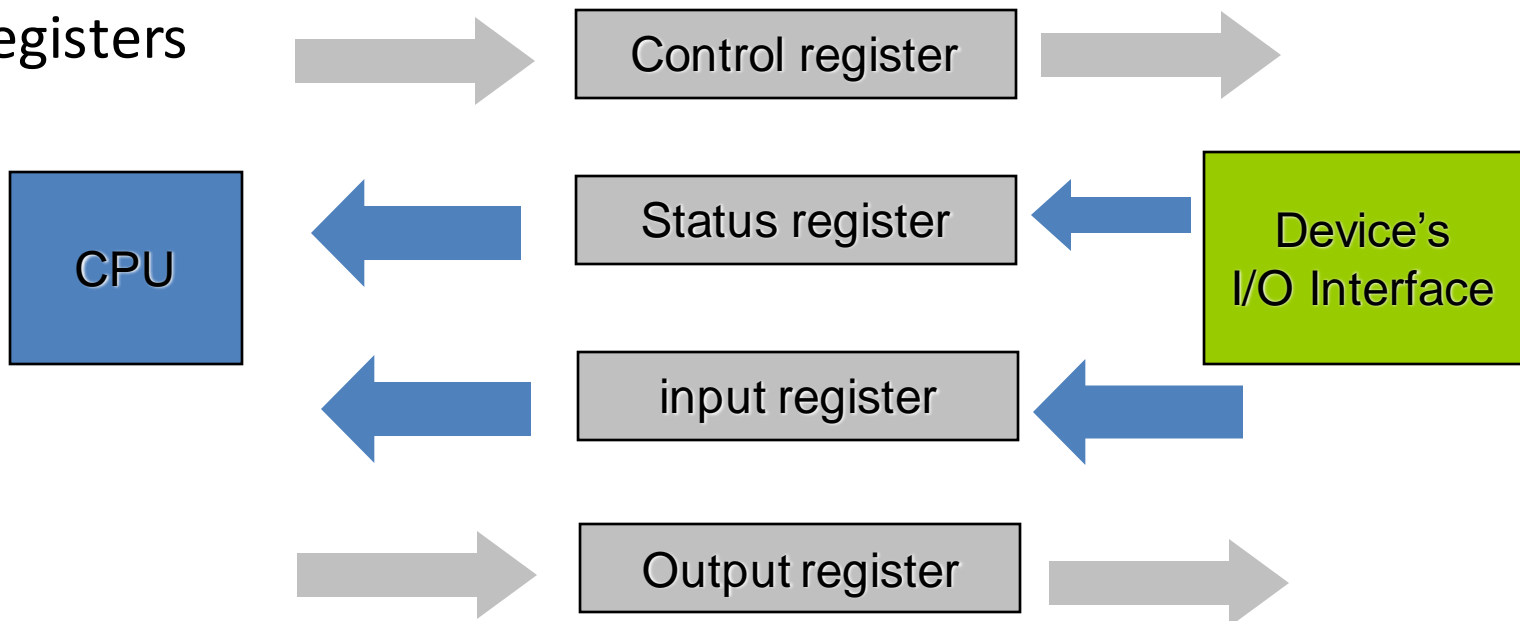
- OS can recover from disk read, device unavailable, transient write failures
- Most return an error number or code when I/O request fails
- System error logs hold problem reports

Improving Performance

- *I/O a major factor in system performance*
 - Demands CPU to execute device driver, kernel I/O code
 - Context switches due to interrupts
 - Data copying
 - Network traffic especially stressful
- Improving Performance
 - Reduce number of *context switches*
 - Reduce *data copying*
 - Reduce *interrupts* by using large transfers, smart controllers, polling
 - Use DMA
 - Balance CPU, memory, bus, and I/O performance for highest throughput

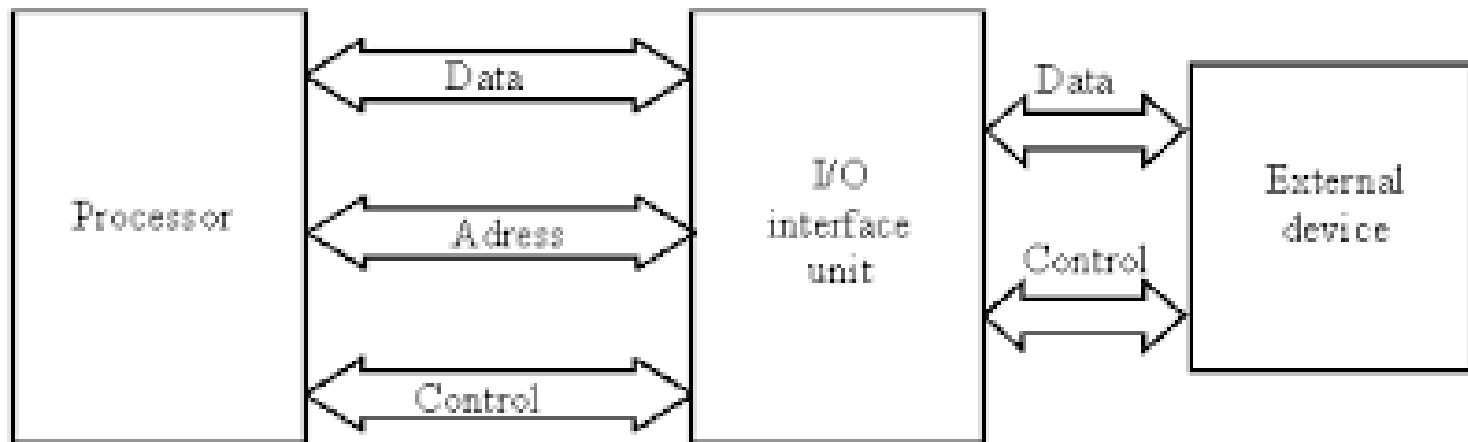
I/O and registers

- The I/O ports of each device are actually structured into a set of specialized registers
- CPU write commands to **control registers**
- CPU reads a value that represents the internal state of the device from **status** register
- CPU fetch data from a device by reading **input** register
- CPU push data to device by writing bytes into **output** registers




What is an I/O interface?

- An I/O interface is a **hardware circuit** inserted between a group of I/O ports and the corresponding device controller.
- It acts as an interpreter that translates the values in the I/O ports into commands and data for the device.



Types of I/O interfaces

- Custom I/O interfaces
 - devoted to one specific hardware device
 - usually, the device controller is on the same card of I/O interfaces.
- General-purpose I/O interface (GPIO)
 - used to connect several different hardware devices



Custom I/O interfaces

- Keyboard interface
- Graphic interface
- Disk interface
- Network interface



- <https://www.coursera.org/lecture/iot-architecture/device-drivers-AL7YG>
- Windows: [Device Driver Development Tutorial – YouTube](#)

<https://www.youtube.com/watch?v=jC0B2kSyKAI>

- Linux: [Learning Linux Device Drivers Development : Find and Create Network Drivers | packtpub.com - YouTube](#)

https://www.youtube.com/watch?v=4Njw_VJUHz4

- <https://github.com/PacktPublishing/Linux-Device-Drivers-Development>